

Machine Learning 1


Angela Abraham

2025-10-21

First up kmeans()

Demo of using kmeans() function in base R. First make up some data with a known structure

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



Now we have some made up data in ‘x’ let’s see how kmeans works with this data

```
k<-kmeans(x, centers=2,nstart=20 )
k
```

Q. How many points are in each cluster

k\$size

```
## [1] 30 30
```

Q. How do we get to the cluster membership / assignment.

k\$cluster


Q. What about cluster centers?

k\$centers

```
##           x         y
## 1  2.644404 -2.978722
## 2 -2.978722  2.644404
```

Now we got to the main results. Let's use them to plot our data with the kmeans result

```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15)
```



Now for Hierarchical Clustering

We will cluster the same data ‘x’ with the ‘`hclust()`’. In this case ‘`hclust()`’ requires a distance matrix as input.


```
hc<- hclust(dist(x))
hc


##
## Call:
## hclust(d = dist(x))
##
## Cluster method : complete
## Distance       : euclidean
## Number of objects: 60
```

Let's plot our hc list result

```
plot(hc)
```

Cluster Dendrogram





#Principal Component Analysis (PCA)

PCA of UK food data


Read data from website and try a few visualizations.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```


	England	Wales	Scotland	N.Ireland
## Cheese	105	103	103	66
## Carcass_meat	245	227	242	267
## Other_meat	685	803	750	586
## Fish	147	160	122	93
## Fats_and_oils	193	235	184	209
## Sugars	156	175	147	139
## Fresh_potatoes	720	874	566	1033
## Fresh_Veg	253	265	171	143
## Other_Veg	488	570	418	355
## Processed_potatoes	198	203	220	187
## Processed_Veg	360	365	337	334
## Fresh_fruit	1102	1137	957	674
## Cereals	1472	1582	1462	1494
## Beverages	57	73	53	47

```
## Soft_drinks      1374  1256   1572   1506
## Alcoholic_drinks 375   475    458    135
## Confectionery     54    64     62     41
```


```
cols<-rainbow(nrow(x))
barplot(as.matrix(x), col=cols )
```



```
barplot(as.matrix(x), col=cols, beside=TRUE)
```



```
pairs(x, col=cols)
```



PCA to the rescue!! The main base R PCA function is called ‘prcomp()’ and we will need to give it to the transpose of our input data!

```
pca<- prcomp(t(x))
```

There is a nice summary of how well PCA is doing

```
summary(pca)
```


```
## Importance of components:
##                 PC1        PC2        PC3        PC4
## Standard deviation   324.1502  212.7478  73.87622 3.176e-14
## Proportion of Variance 0.6744    0.2905   0.03503 0.000e+00
## Cumulative Proportion 0.6744    0.9650   1.00000 1.000e+00
```

```
attributes(pca)
```

```
## $names
## [1] "sdev"      "rotation"   "center"     "scale"      "x"
##
## $class
## [1] "prcomp"
```


To make our new PCA plot (a.k.a PCA score plot) we access ‘pca\$x’

```
plot(pca$x[,1], pca$x[,2])
text(pca$x[,1], pca$x[,2], colnames(x))
```



color up the plot

```
country_cols<- c("orange", "red", "blue", "green")
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
text(pca$x[,1], pca$x[,2], col=country_cols)
```



PCA of RNA-Seq data

Read in data from website

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2 wt3 wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458 408 429 420 90 88 86 90 93
## gene2 219 200 204 210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4 783 792 829 856 760 849 856 835 885 894
## gene5 181 249 204 244 225 277 305 272 270 279
## gene6 460 502 491 491 493 612 594 577 618 638
```


```
pca<- prcomp(t(rna.data))
summary(pca)
```

```
## Importance of components:
##                               PC1       PC2       PC3       PC4       PC5       PC6
## Standard deviation    2214.2633 88.9209 84.33908 77.74094 69.66341 67.78516
## Proportion of Variance 0.9917 0.0016 0.00144 0.00122 0.00098 0.00093
## Cumulative Proportion 0.9917 0.9933 0.99471 0.99593 0.99691 0.99784
```

```
##          PC7      PC8      PC9      PC10
## Standard deviation   65.29428 59.90981 53.20803 2.684e-13
## Proportion of Variance 0.00086  0.00073  0.00057  0.000e+00
## Cumulative Proportion 0.99870  0.99943  1.00000  1.000e+00
```

Do our PCA plot of this RNA-Seq data

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
text(pca$x[,1], pca$x[,2], colnames(rna.data))
```

